

Technical Assessment: Intelligent Library System ("LuminaLib")

Role: Senior Backend & GenAI Engineer

Stack: Python (FastAPI), PostgreSQL, Generative AI (Llama 3), Docker

1. Project Context

You are architecting "LuminaLib," a next-generation library system. This is not a simple metadata CRUD application. It is a content-aware system that manages actual book files, uses a local LLM to synthesize reader sentiment, and employs machine learning to drive discovery.

Evaluation Goal: We are looking for production-grade engineering. We value **Clean Architecture**, **SOLID principles**, **Configurability**, and **Async mastery** over quick-and-dirty scripts.

2. Functional Requirements

A. Authentication & User Management

- **Security:** Implement a robust authentication system. The API must not be open; it requires **JWT-based stateless authentication**.
- **User Lifecycle:** Provide endpoints for **Signup**, **Login**, **Profile Management**, and **Signout**.
- **Design Challenge:** You must design the users table and associated security flows (password hashing, token expiration) from scratch.

B. Book Ingestion & Management (The "Content" Layer)

- **True Ingestion:** Creating a book is not just inserting a title and author into a DB. You must implement a mechanism to **upload and store the actual book content** (text or PDF).
 - *Storage:* You may simulate S3 using a local container (like MinIO) or manage local file storage, but the architecture must treat storage as an abstraction.
- **Full CRUD:** Implement Create, Read, Update, and Delete operations for book metadata and content.
- **Library Mechanics:** Implement **Borrow** and **Return** functionality.
 - *Constraint:* Users cannot review a book they haven't borrowed.

C. The Intelligence Layer (GenAI & ML)

- **Book Summarization:** Upon ingestion, the system must asynchronously read the book content and generate a summary using **LLM** of your choice.
- **Review Analysis:** When a user submits a review (POST /books/{id}/reviews), it must trigger a background task. This task uses the LLM to update a "rolling consensus" or summary of what users feel about the book.
- **Recommendation Engine:**
 - **User Preferences:** The requirement is vague on purpose. You must design a DB schema to store "User Preferences." How you model this (explicit tags, implicit borrow history, etc.) is part of the test.
 - **ML Model:** Implement a recommendation algorithm (e.g., Content-Based or Collaborative Filtering) that uses these preferences to suggest books.

3. Architecture & Code Quality

- **Clean Architecture:** Use **Dependency Injection (DI)** and **Interface-Driven Development**.
 - *Extensibility:* We will evaluate how easily we can swap components. For example, can we switch the file storage from Local Disk to AWS S3 by changing a single config line? Can we swap Llama 3 for OpenAI without rewriting business logic?
- **Code Standards:**
 - **Linting & Formatting**
 - **Proper import sequencing**
 - **Type Safety**

4. Deployment & Infrastructure

Shutterstock

Explore

- **Dockerization:** The application must be fully containerized.
- **One-Command Start:** We expect a docker-compose.yml file. We must be able to run a single command (e.g., docker-compose up --build) to spin up:
 1. The API Service
 2. The PostgreSQL Database
 3. The Local LLM Service (e.g., via Ollama or a mock container if hardware is limited)

- **Configuration:** All environment variables (DB URLs, API Keys, Model paths) must be managed via .env files and injected into the containers.

5. API Specification (Expanded)

Your RESTful API must support at least the following:

Domain	Method	Endpoint	Description
Auth	POST	/auth/signup	Register a new user.
	POST	/auth/login	Return JWT access token.
Books	POST	/books	Upload book file & metadata. Triggers async summary.
	GET	/books	List books (pagination required).
	PUT	/books/{id}	Update book details.
	DELETE	/books/{id}	Remove book and associated file.
	POST	/books/{id}/borrow	User borrows a book.
	POST	/books/{id}/return	User returns a book.
Reviews	POST	/books/{id}/reviews	Submit review. Triggers async sentiment analysis.
Intel	GET	/books/{id}/analysis	Get GenAI-aggregated summary of all reviews.
	GET	/recommendations	Get ML-based suggestions for the current user.

6. Deliverables

1. **Source Code:** A clean, structured repository.
2. **docker-compose.yml:** The orchestration file.
3. **ARCHITECTURE.md:** A document explaining:
 - o Why you chose your specific DB Schema for User Preferences.
 - o How you handled the "Async" aspect of LLM generation (Queues? Background Tasks?).
 - o Your strategy for the ML recommendation model.
4. **README.md:** Instructions to run the application using the one-command start.

Evaluation Rubric

1. **Modularity:** Can we swap the Storage/LLM providers easily?
2. **Docker Proficiency:** Does the compose file work seamlessly?
3. **Code Hygiene:** Are imports sorted? Is the code linted?

4. **GenAI Implementation:** Is the prompt engineering handled in a structured, reusable way?